



Keeping your GPU fed without getting bitten

Tobias Hector
May 2017


Introduction

- You have delicious draw calls
 - Yummy!




Introduction

- You have delicious draw calls
 - Yummy!
- Your GPU wants to eat them
 - It's really hungry



Introduction

- You have delicious draw calls
 - Yummy!
- Your GPU wants to eat them
 - It's really hungry
- Keep it fed at all times
 - So it keeps making pixels



Introduction

- You have delicious draw calls
 - Yummy!
- Your GPU wants to eat them
 - It's really hungry
- Keep it fed at all times
 - So it keeps making pixels
- Don't want it biting your hand
 - Look at those teeth!



Keeping it fed

- GPU needs a constant supply of food
 - It doesn't want to wait
- Certain foods are tough to digest
 - Provide multiple operations to hide stalls
- Draw calls provide a variety of nutrition
 - Vertex work, raster work, tessellation, vitamins A-K, etc.

Keeping it fed

System				
CPU	0		1	
GPU		0		1

Keeping it fed

System				
CPU	0	1	2	
GPU		0	1	2

Keeping it fed

GPU				
Vertex	0		1	
Fragment		0		1

Keeping it fed

GPU				
Vertex	0	1	2	
Fragment		0	1	2

Not getting bitten


- GPU eating from lots of different plates
 - Don't touch anything it's using!
- It doesn't want a mouthful of beef choc chip ice cream
 - Don't change data whilst it's accessing a resource
- Hey I'm eating that!
 - Don't delete resources whilst the GPU is still using them



Tear Point #1 --->

Tear Point #2 --->








On to the serious bits...

Terminology

- **Operation**
 - An executable task
- **Execution Dependency**
 - Guarantee for one set of operations to wait on another set of operations
- **Memory Space**
 - RAM, caches or registers
- **Write Propagation**
 - Operation that copies a written value between memory spaces
- **Memory Dependency**
 - Execution dependency including write propagations

Vulkan Memory Spaces



Vulkan Memory Spaces

“VRAM”

Device
Memory

“Caches”

Host
Transfer

Pipeline
Stages

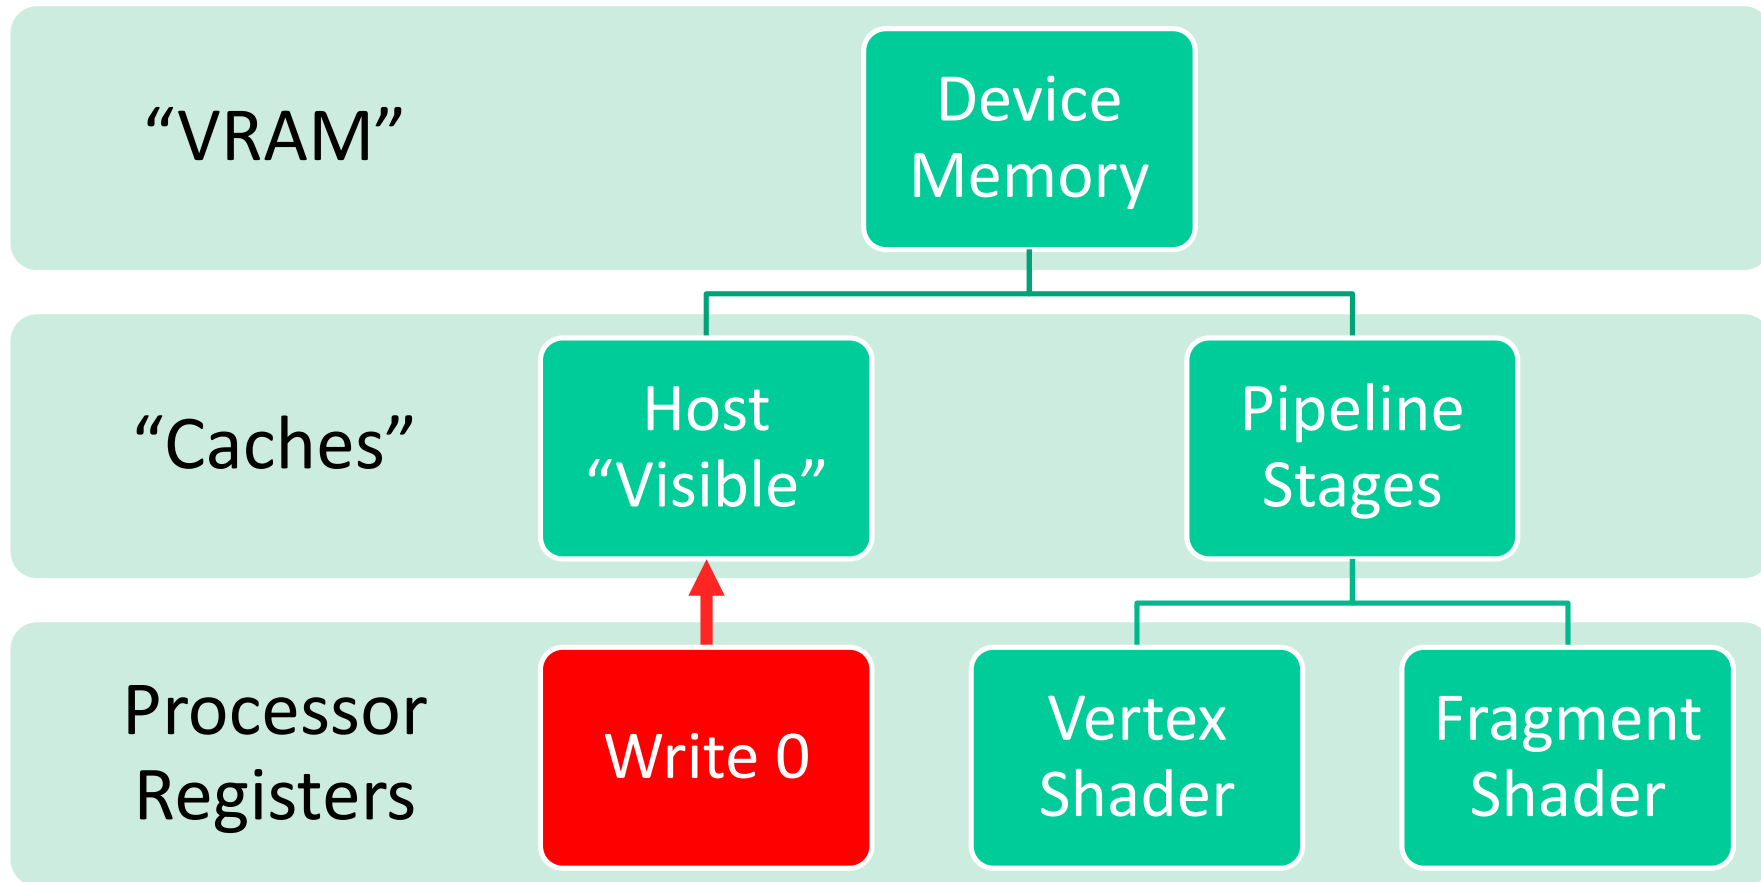
Processor
Registers

CPU

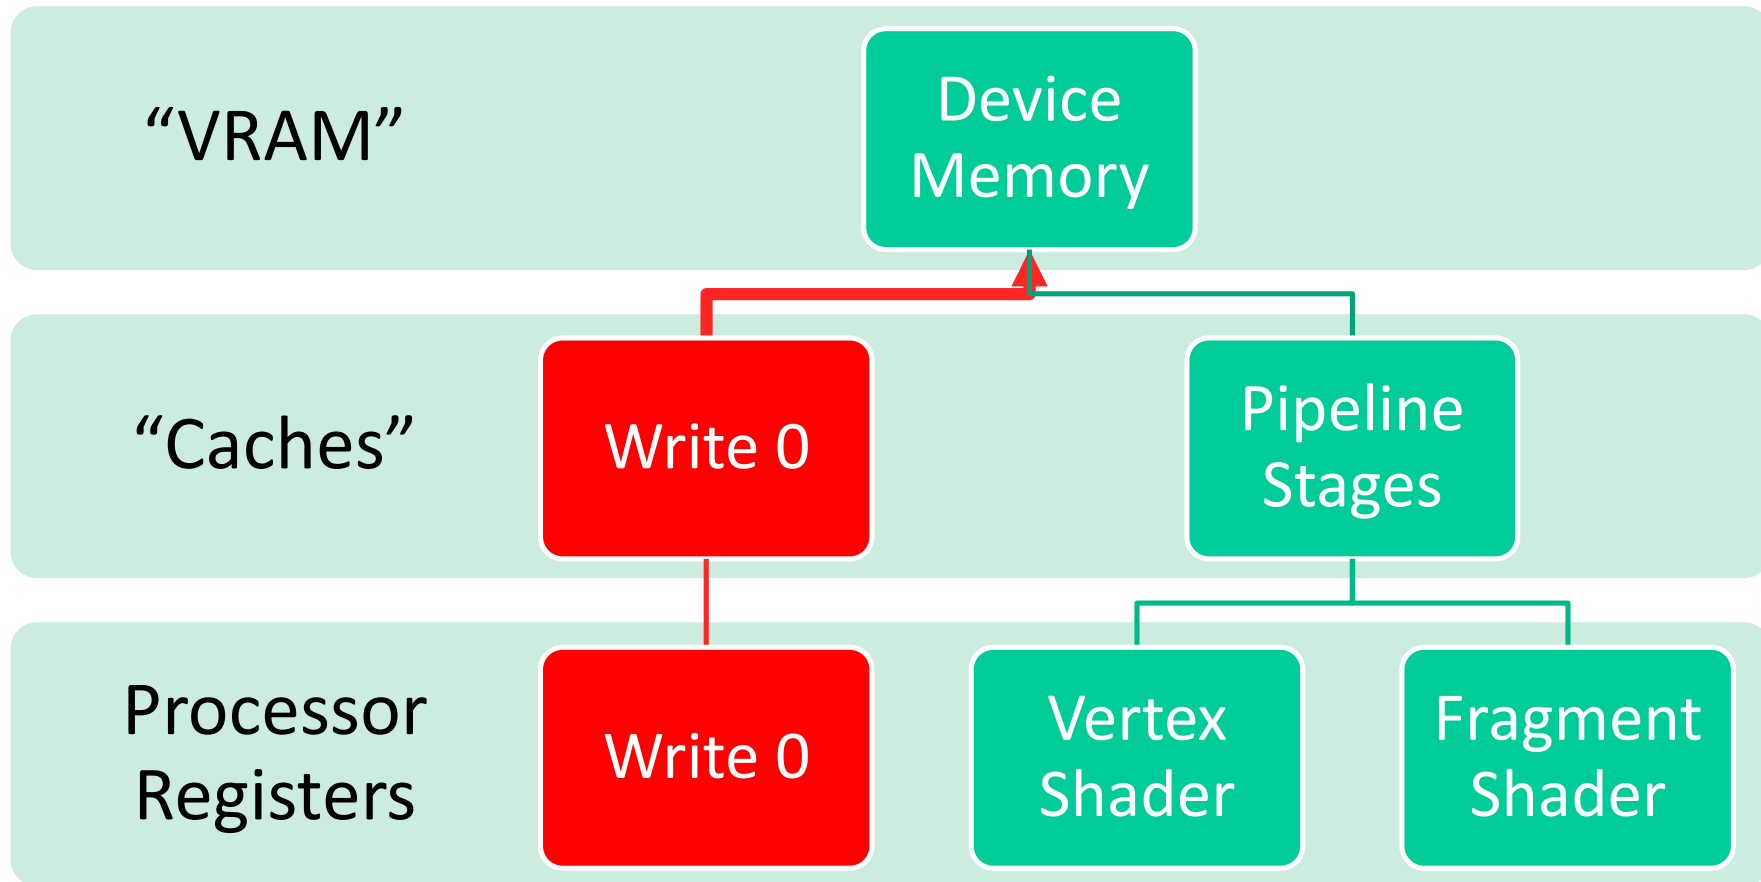
Vertex
Shader

Fragment
Shader

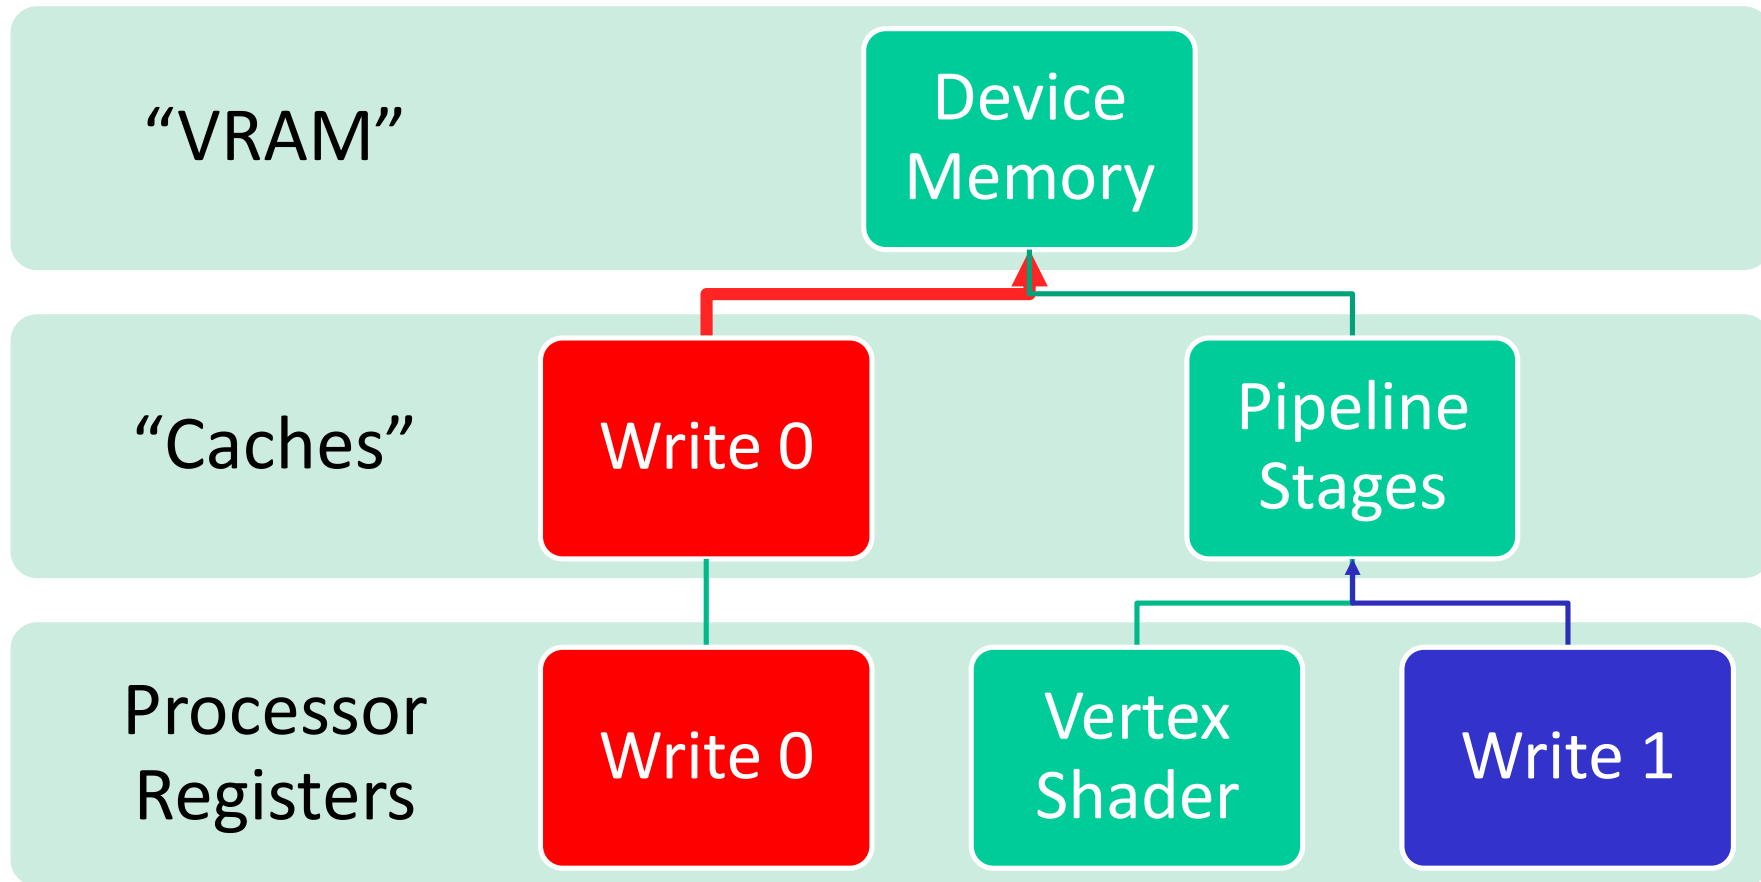
Vulkan Memory Spaces - Data Hazards



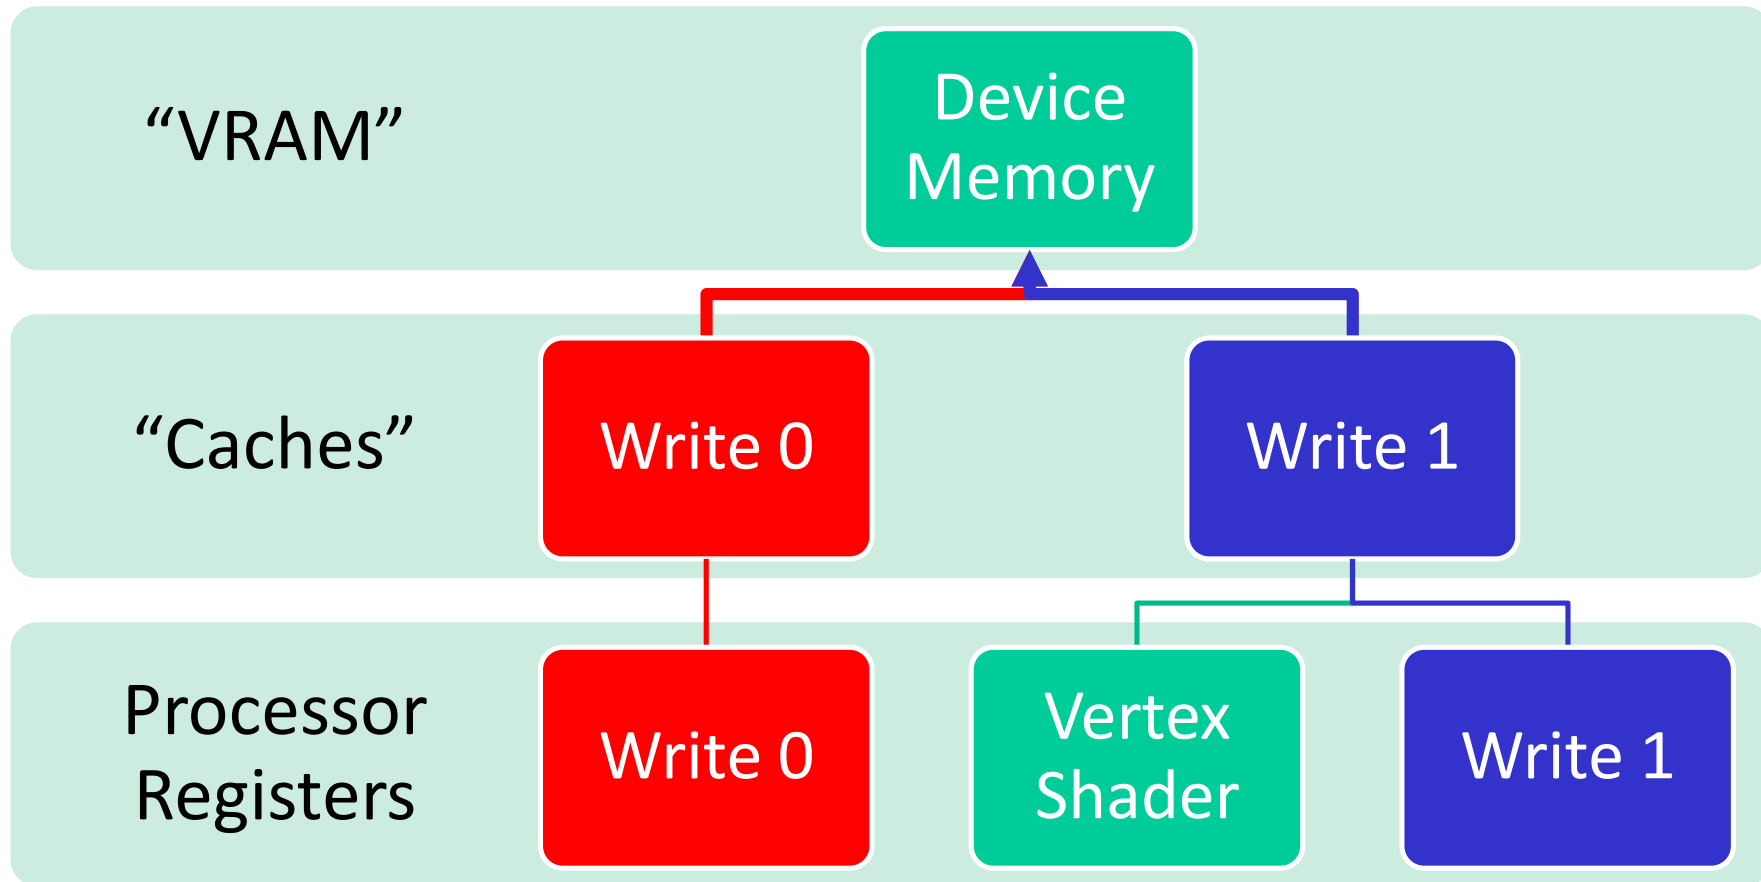
Vulkan Memory Spaces - Data Hazards



Vulkan Memory Spaces - Data Hazards



Vulkan Memory Spaces - Data Hazards



Vulkan Memory Spaces - Data Hazards

“VRAM”

????

“Caches”

Write 0

Write 1

Processor
Registers

Write 0

Vertex
Shader

Write 1

Vulkan Memory Spaces - Data Hazards

“VRAM”

????

“Caches”

????

????

Processor
Registers

????

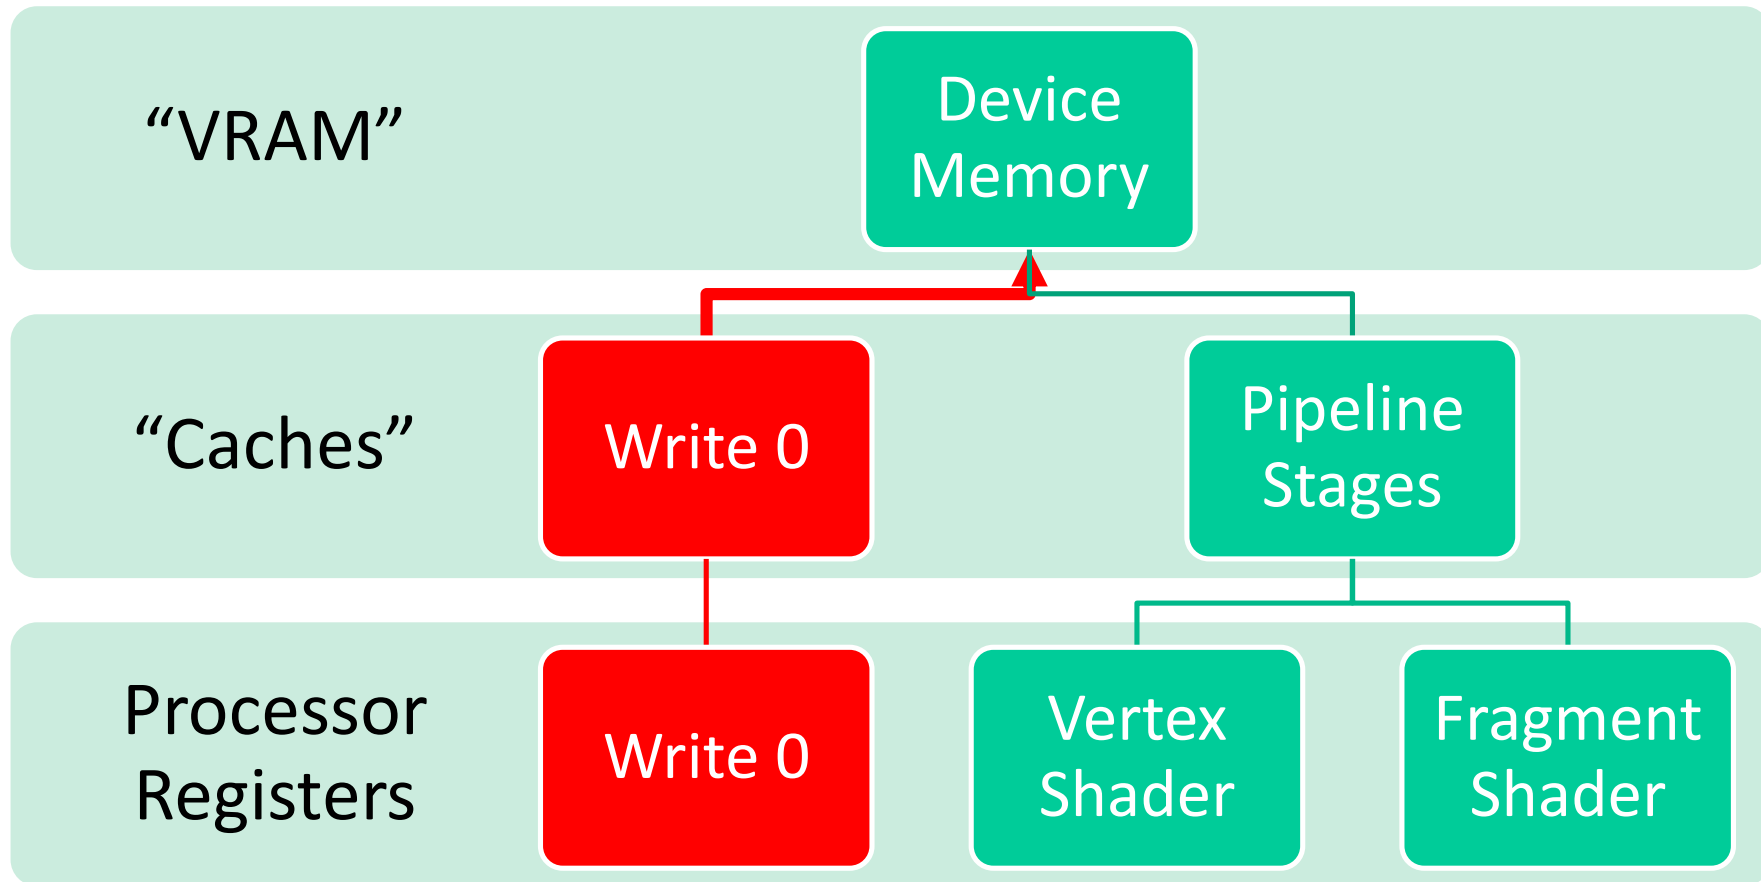
????

????

Oops

- Let's try that again...
- This time, with a memory dependency!

Vulkan Memory Spaces - Data Hazards



Vulkan Memory Spaces - Data Hazards

“VRAM”

Write 0

“Caches”

Write 0

Write 0

Processor
Registers

Write 0

Vertex
Shader

Write 0

Vulkan Memory Spaces - Data Hazards

“VRAM”

Write 0

“Caches”

Write 0

Write 0

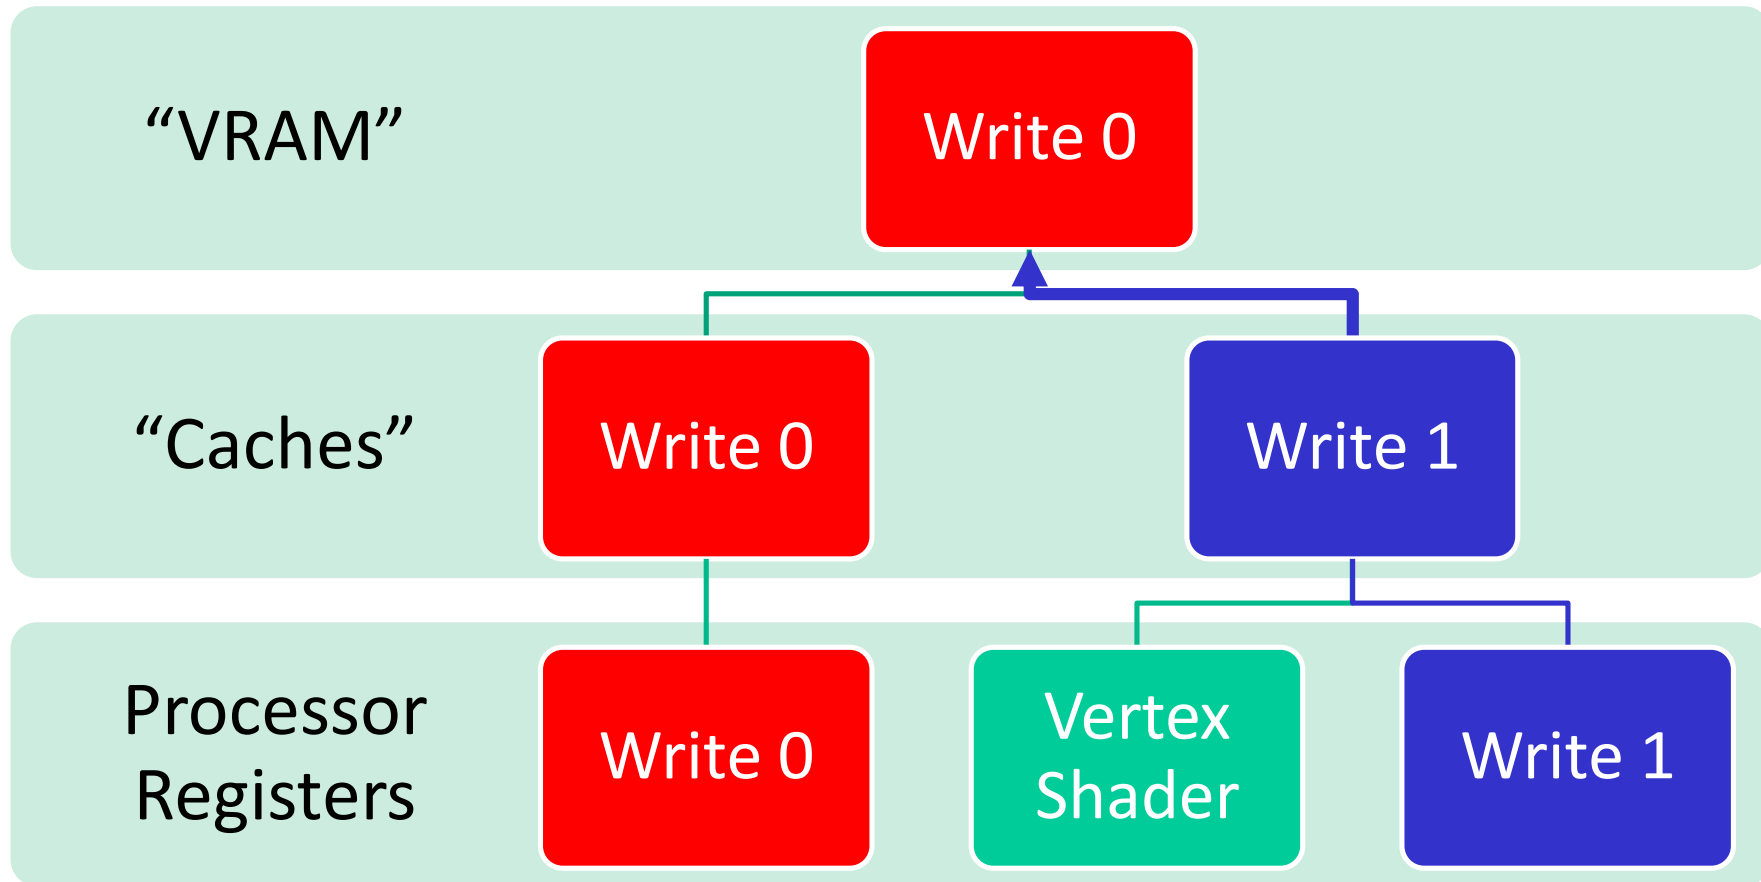
Processor
Registers

Write 0

Vertex
Shader

Write 1

Vulkan Memory Spaces - Data Hazards



Vulkan Memory Spaces - Data Hazards

“VRAM”

Write 1

“Caches”

Write 1

Write 1

Processor
Registers

Write 1

Vertex
Shader

Write 1

How do I do that?

- Synchronization primitives!

Synchronization Types

- 3 types of explicit synchronization in Vulkan
 - Pipeline Barriers, Events and Subpass Dependencies
 - Within a queue
 - Explicit memory dependencies
 - Semaphores
 - Between Queues
 - Fences
 - Whole queue operations to CPU

OpenGL has just two, very coarse synchronization primitives: memory barriers and fences. They are loosely similar to the equivalently named concepts in Vulkan

Pipeline Barriers

- Pipeline Barriers
 - Precise set of pipeline stages
 - Memory Barriers to execute
 - Single point in time

Executing a pipeline barrier is similar to a `glMemoryBarrier` call, though with much more control.

```
void vkCmdPipelineBarrier(  
    VkCommandBuffer          commandBuffer,  
    VkPipelineStageFlags     srcStageMask,  
    VkPipelineStageFlags     dstStageMask,  
    VkDependencyFlags         dependencyFlags,  
    uint32_t                 memoryBarrierCount,  
    const VkMemoryBarrier*   pMemoryBarriers,  
    uint32_t                 bufferMemoryBarrierCount,  
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,  
    uint32_t                 imageMemoryBarrierCount,  
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```

Events

- Events
 - Same info as Pipeline Barriers
 - ...but operate over a range

```
void vkCmdSetEvent(
    VkCommandBuffer      commandBuffer,
    VkEvent               event,
    VkPipelineStageFlags stageMask);

void vkCmdResetEvent(
    VkCommandBuffer      commandBuffer,
    VkEvent               event,
    VkPipelineStageFlags stageMask);

void vkCmdWaitEvents(
    VkCommandBuffer      commandBuffer,
    uint32_t             eventCount,
    const VkEvent*       pEvents,
    VkPipelineStageFlags srcStageMask,
    VkPipelineStageFlags dstStageMask,
    uint32_t             memoryBarrierCount,
    const VkMemoryBarrier* pMemoryBarriers,
    uint32_t             bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,
    uint32_t             imageMemoryBarrierCount,
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```

Events

- **Events**
 - Same info as Pipeline Barriers
 - ...but operate over a range
- **CPU interaction**
 - No explicit CPU wait

```
VkResult vkSetEvent(  
    VkDevice      device,  
    VkEvent       event);
```

```
VkResult vkResetEvent(  
    VkDevice      device,  
    VkEvent       event);
```

```
VkResult vkGetEventStatus(  
    VkDevice      device,  
    VkEvent       event);
```

Events

- **Events**
 - Same info as Pipeline Barriers
 - ...but operate over a range
- **CPU interaction**
 - No explicit CPU wait
- **Warning!**
 - May timeout
 - Set events soon after submission
 - Could you just defer submission?

```
VkResult vkSetEvent(  
    VkDevice device,  
    VkEvent event);
```

```
VkResult vkResetEvent(  
    VkDevice device,  
    VkEvent event);
```

```
VkResult vkGetEventStatus(  
    VkDevice device,  
    VkEvent event);
```



Pipeline Barriers vs Events

- **Use pipeline barriers for point synchronization**
 - Dependant operation immediately precedes operation that depends on it
 - May be more optimal than set/wait event pair
- **Use events if other work possible between two operations**
 - Set immediately after the dependant operation
 - Wait immediately before the operation that depends on it
 - Allows more overlap of work
- **Use events for CPU/GPU synchronization**
 - Memory accesses between processors
 - Late latching of data to reduce latency

Memory Barriers

- Defines write propagations
 - Between “visible” and device memory
- Three types...

OpenGL’s memory barriers imply execution dependencies, which Vulkan memory barriers do not - execution dependencies are provided by a pipeline barrier, event or subpass dependency.

Global Memory Barriers

- Global Memory Barriers
 - All memory used by **access types**
 - Flushes/invalidates whole caches
- Use when many resources transition
 - Cheaper than one-by-one
 - Don't transition unnecessarily!
- User defines prior access
 - Driver not tracking for you

```
typedef struct VkMemoryBarrier {  
    VkStructureType      sType;  
    const void*          pNext;  
    VkAccessFlags        srcAccessMask;  
    VkAccessFlags        dstAccessMask;  
} VkMemoryBarrier;
```

Buffer Barriers

- Buffer Barriers
 - A single **buffer range**
 - Defines **access types**
 - Defines **queue ownership**
- Buffer Range
 - Offset and size within a buffer
- Queue Ownership
 - Defines which queue families are *allowed* to access a write

```
typedef struct VkBufferMemoryBarrier {  
    VkStructureType      sType;  
    const void*          pNext;  
    VkAccessFlags         srcAccessMask;  
    VkAccessFlags         dstAccessMask;  
    uint32_t             srcQueueFamilyIndex;  
    uint32_t             dstQueueFamilyIndex;  
    VkBuffer              buffer;  
    VkDeviceSize          offset;  
    VkDeviceSize          size;  
} VkBufferMemoryBarrier;
```

Image Barriers

- Image Barriers

- A single **image subresource range**
- Defines **access types**
- Defines **queue ownership**
- Defines **image layout**

- Image subresource range

- Specific levels/layers of an image

- Image layouts

- Additional access information for images
- Enables GPU image compression
- Use GENERAL rather than frequent switching

```
typedef struct VkImageMemoryBarrier {  
    VkStructureType      sType;  
    const void*          pNext;  
    VkAccessFlags         srcAccessMask;  
    VkAccessFlags         dstAccessMask;  
    VkImageLayout         oldLayout;  
    VkImageLayout         newLayout;  
    uint32_t              srcQueueFamilyIndex;  
    uint32_t              dstQueueFamilyIndex;  
    VkImage               image;  
    VkImageSubresourceRange subresourceRange;  
} VkImageMemoryBarrier;
```

Example - Texture Upload

```
// Read image from file, flush to 'host visible' memory space
fread(mappedBufferMemory, 1, imageDataSize, imageFile);
vkFlushMappedMemoryRanges(..., {mappedBufferMemory, ...});
// Transition the buffer from host write to transfer read
bBarrier.srcAccessMask = VK_ACCESS_HOST_WRITE_BIT; // Buffer being written to
bBarrier.dstAccessMask = VK_ACCESS_TRANSFER_READ_BIT; // Buffer will be read
// Transition the image to transfer destination
iBarrier.srcAccessMask = 0; // No prior access
iBarrier.dstAccessMask = VK_ACCESS_TRANSFER_WRITE_BIT; // Get image prepared to be transferred to
iBarrier.oldLayout = VK_IMAGE_LAYOUT_UNDEFINED; // No prior access
iBarrier.newLayout = VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL; // Get image prepared to be transferred to
// Pipeline barrier for pre-transfer memory dependency - buffer write was on the host
vkCmdPipelineBarrier(commandBuffer, VK_PIPELINE_STAGE_HOST_BIT, VK_PIPELINE_STAGE_TRANSFER_BIT, &bBarrier, &iBarrier);
// Copy from buffer to image
vkCmdCopyBufferToImage(commandBuffer, srcBuffer, image, VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL, 1, &copy);
// Transition the image from transfer destination to shader read
iBarrier.srcAccessMask = VK_ACCESS_TRANSFER_WRITE_BIT; // Image was just written to
iBarrier.dstAccessMask = VK_ACCESS_SHADER_READ_BIT; // Get image prepared to be read by a shader
iBarrier.oldLayout = VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL; // Image was just written to
iBarrier.newLayout = VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL; // Get image prepared to be read by a shader
// Pipeline barrier for post-transfer memory dependency - fragment shader will read
vkCmdPipelineBarrier(commandBuffer, VK_PIPELINE_STAGE_TRANSFER_BIT, VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT, &iBarrier);
```

Usually though...

- **Many pipeline barrier operations doable via a render pass**
 - At least for drawing operations
- **Render passes are a dependency graph**
 - Allows driver to plan ahead of time how to execute
 - Typically much more efficient than manual barriers

Subpass Dependencies

- Subpass dependencies
 - Similar info to Pipeline Barriers
 - Explicitly between **two subpasses**
- Memory barriers
 - Implicit for attachments
 - **Explicit for other resources**
- Framebuffer-local dependencies
 - Same fragment/sample location
 - Cheap for most implementations
 - Use region dependency flag:
 - VK_DEPENDENCY_BY_REGION_BIT

```
typedef struct VkSubpassDependency {  
    uint32_t                srcSubpass;  
    uint32_t                dstSubpass;  
    VkPipelineStageFlags    srcStageMask;  
    VkPipelineStageFlags    dstStageMask;  
    VkAccessFlags           srcAccessMask;  
    VkAccessFlags           dstAccessMask;  
    VkDependencyFlags       dependencyFlags;  
} VkSubpassDependency;
```

Subpass Self-Dependencies

- Subpass self-dependencies
 - Subpasses can wait on themselves
 - A pipeline barrier in the subpass
- Forward progress only
 - Can't wait on later stages
 - Must wait on earlier or same stage
- Only framebuffer-local for fragments
 - Must use flag:
 - VK_DEPENDENCY_BY_REGION_BIT

```
typedef struct VkSubpassDependency {
    uint32_t                srcSubpass;
    uint32_t                dstSubpass;
    VkPipelineStageFlags    srcStageMask;
    VkPipelineStageFlags    dstStageMask;
    VkAccessFlags           srcAccessMask;
    VkAccessFlags           dstAccessMask;
    VkDependencyFlags        dependencyFlags;
} VkSubpassDependency;

void vkCmdPipelineBarrier(
    VkCommandBuffer          commandBuffer,
    VkPipelineStageFlags     srcStageMask,
    VkPipelineStageFlags     dstStageMask,
    VkDependencyFlags        dependencyFlags,
    uint32_t                memoryBarrierCount,
    const VkMemoryBarrier*   pMemoryBarriers,
    uint32_t                bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,
    uint32_t                imageMemoryBarrierCount,
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```

Subpass External Dependencies

- Subpass external dependencies
 - Wait on 'external' operations
 - vkCmdWaitEvent in the subpass
 - Events set outside the render pass
- Very useful for common dependencies
 - Use to move between PRESENT_SRC and COLOR_ATTACHMENT_OUTPUT
 - Avoids need for pipeline barriers

```
typedef struct VkSubpassDependency {
    uint32_t                srcSubpass;
    uint32_t                dstSubpass;
    VkPipelineStageFlags    srcStageMask;
    VkPipelineStageFlags    dstStageMask;
    VkAccessFlags           srcAccessMask;
    VkAccessFlags           dstAccessMask;
    VkDependencyFlags       dependencyFlags;
} VkSubpassDependency;

void vkCmdWaitEvents(
    VkCommandBuffer          commandBuffer,
    uint32_t                eventCount,
    const VkEvent*           pEvents,
    VkPipelineStageFlags     srcStageMask,
    VkPipelineStageFlags     dstStageMask,
    uint32_t                memoryBarrierCount,
    const VkMemoryBarrier*   pMemoryBarriers,
    uint32_t                bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,
    uint32_t                imageMemoryBarrierCount,
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```


Example - Acquire, Render, Present

```
// Subpass dependency to express that an attachment has just been acquired
acquiredDependency.srcSubpass = VK_SUBPASS_EXTERNAL;
acquiredDependency.dstSubpass = 0; // First subpass it's used in
acquiredDependency.srcStageMask = VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT; // Semaphore/Submit guarantees this is sufficient
acquiredDependency.dstStageMask = VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT; // Latest possible stage
acquiredDependency.srcAccessMask = 0; // Previous semaphore will guarantee it is in device memory
acquiredDependency.dstAccessMask = VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT; // Access type
```

```
// Subpass dependency to express that an attachment will be presented
presentDependency.srcSubpass = 4; // Last subpass it's used in
presentDependency.dstSubpass = VK_SUBPASS_EXTERNAL;
presentDependency.srcStageMask = VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT; // Earliest possible stage
presentDependency.dstStageMask = VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT; // Semaphore/Submit guarantees this is sufficient
presentDependency.srcAccessMask = VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT; // Access type
presentDependency.dstAccessMask = 0; // Previous semaphore will guarantee it is in device memory
```

```
// Acquire image, render to it, then present it (render loop!)
vkAcquireImageKHR(...);
vkQueueSubmit(...); // Includes execution of render pass
vkQueuePresentKHR(...);
```

Semaphores

- Semaphores
 - Used to synchronize queues
 - Not necessary for single-queue
- Fairly coarse
 - Per submission batch
 - E.g. a set of command buffers
 - Multiple per submit command
- Some implicit memory dependencies
 - Writes propagated between all device “visible” memory spaces
 - Not guaranteed visible to host

```
typedef struct VkSubmitInfo {  
    VkStructureType           sType;  
    const void*               pNext;  
    uint32_t                  waitSemaphoreCount;  
    const VkSemaphore*        pWaitSemaphores;  
    const VkPipelineStageFlags* pWaitDstStageMask;  
    uint32_t                  commandBufferCount;  
    const VkCommandBuffer*     pCommandBuffers;  
    uint32_t                  signalSemaphoreCount;  
    const VkSemaphore*        pSignalSemaphores;  
} VkSubmitInfo;
```

Example - Acquire, Render, Present 2

```
// Acquire an image. Pass in a semaphore to be signalled
```

```
vkAcquireNextImageKHR(device, swapchain, UINT64_MAX, acquireSemaphore, VK_NULL_HANDLE, &imageIndex);
```

```
// Submit command buffers
```

```
submitInfo.waitSemaphoreCount = 1;
```

```
submitInfo.pWaitSemaphores = &acquireSemaphore;
```

```
submitInfo.commandBufferCount = 1;
```

```
submitInfo.pCommandBuffers = &commandBuffer;
```

```
submitInfo.signalSemaphoreCount = 1;
```

```
submitInfo.pSignalSemaphores = &graphicsSemaphore;
```

```
vkQueueSubmit(graphicsQueue, 1, &submitInfo, fence);
```

```
// Present images to the display
```

```
presentInfo.waitSemaphoreCount = 1;
```

```
presentInfo.pWaitSemaphores = &graphicsSemaphore;
```

```
presentInfo.swapchainCount = 1;
```

```
presentInfo.pSwapchains = &swapchain;
```

```
presentInfo.pImageIndices = &imageIndex;
```

```
vkQueuePresentKHR(presentQueue, &presentInfo);
```

Fences

- **Fences**
 - Used to synchronize queue to CPU
- **Very coarse grain**
 - Per queue submit command
- **Implicit memory dependency**
 - Writes propagated between all device “visible” memory spaces
 - Not guaranteed visible to host

GL's fences are like a combination of a semaphore and a fence in Vulkan - they can synchronize GPU and CPU in multiple ways at a coarse granularity.

```
VkResult vkQueueSubmit(  
    VkQueue          queue,  
    uint32_t         submitCount,  
    const VkSubmitInfo* pSubmits,  
    VkFence          fence);  
  
VkResult vkResetFences(  
    VkDevice          device,  
    uint32_t          fenceCount,  
    const VkFence*    pFences);  
  
VkResult vkGetFenceStatus(  
    VkDevice          device,  
    VkFence           fence);  
  
VkResult vkWaitForFences(  
    VkDevice          device,  
    uint32_t          fenceCount,  
    const VkFence*    pFences,  
    VkBool32          waitAll,  
    uint64_t           timeout);
```

Other important synchronization...

- Implicit... ish?

Queue Submit

- **Queue Submission**
 - Used to push operations to a device
- **Guaranteed forward progress**
 - Operations will execute once pushed
- **Implicit memory dependency**
 - Writes propagated from 'host visible' to all 'device visible' memory spaces

```
VkResult vkQueueSubmit(  
    VkQueue          queue,  
    uint32_t         submitCount,  
    const VkSubmitInfo* pSubmits,  
    VkFence          fence);
```

Wait Idle

- Ensures execution completes
 - VERY heavy-weight
- **vkQueueWaitIdle**
 - Wait for queue operations to finish
 - Equivalent to waiting on a fence
- **vkDeviceWaitIdle**
 - Waits for device operations to finish
 - Includes vkQueueWaitIdle for queues

```
VkResult vkQueueWaitIdle(  
    VkQueue queue);  
  
VkResult vkDeviceWaitIdle(  
    VkDevice device);
```

Wait Idle

- Ensures execution completes
 - VERY heavy-weight
- **vkQueueWaitIdle**
 - Wait for queue operations to finish
 - Equivalent to waiting on a fence
- **vkDeviceWaitIdle**
 - Waits for device operations to finish
 - Includes vkQueueWaitIdle for queues
- **Warning!**
 - Only use for tear-down
 - Will guarantee no overlap

```
VkResult vkQueueWaitIdle(  
    VkQueue queue);  
  
VkResult vkDeviceWaitIdle(  
    VkDevice device);
```



Programmer Guidelines

- **Specify EXACTLY the right amount of synchronization**
 - Too much and you risk starving your GPU
 - Miss any and your GPU will bite you
- **Use the validation layers to help!**
 - Won't catch everything, improving over time
- **Pay particular attention to the pipeline stages**
 - Fiddly but become intuitive as you use them
- **Consider Image Layouts**
 - If your GPU can save bandwidth it will
- **Prefer render passes**
 - Driver able to plan workloads efficiently
- **Pay attention to implicit dependencies**
 - Submit and Semaphores guarantee a lot - don't add more!
- **Different behaviour depending on implementation**
 - Test/Tune on every platform you can find!

Keep your GPU fed without getting bitten!

Questions?



BACKUP SLIDES

Example - Compute to Draw Indirect

```
// Add a subpass dependency to express the wait on an external event
externalDependency.srcSubpass = VK_SUBPASS_EXTERNAL;
externalDependency.srcStageMask = VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT;
externalDependency.dstStageMask = VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT;
externalDependency.srcAccessMask = VK_ACCESS_SHADER_WRITE_BIT;
externalDependency.dstAccessMask = VK_ACCESS_INDIRECT_COMMAND_READ_BIT;
```

```
// Dispatch a compute shader that generates indirect command structures
vkCmdDispatch(...);
// Set an event that can be later waited on (same source stage).
vkCmdSetEvent(commandBuffer, event, VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT);
```

```
vkCmdBeginRenderPass(...);

// Transition the buffer from shader write to indirect command - details match external dependency!
bufferBarrier.srcAccessMask = VK_ACCESS_SHADER_WRITE_BIT;
bufferBarrier.dstAccessMask = VK_ACCESS_INDIRECT_COMMAND_READ_BIT;
bufferBarrier.buffer = indirectBuffer;
vkCmdWaitEvent(commandBuffer, event, VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT, VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT,
&bufferBarrier);

vkCmdDrawIndirect(commandBuffer, indirectBuffer, ...);
```

Example - Multi-buffering

```
// Have enough resources and fences to have one per in-flight-frame, usually the swapchain image count
VkBuffer buffers[swapchainImageCount];
VkFence fence[swapchainImageCount];

// Can use the index from the presentation engine - 1:1 mapping between swapchain images and resources
vkAcquireNextImageKHR(device, swapchain, UINT64_MAX, semaphore, VK_NULL_HANDLE, &nextIndex);

// Make absolutely sure that the work has completed
vkWaitForFences(device, 1, &fence[nextIndex], true, UINT64_MAX);

// Reset the fences we waited on, so they can be re-used
vkResetFences(device, 1, &fence[nextIndex]);

// Change the data in your per-frame resources (with appropriate events/barriers!)
...

// Submit any work to the queue, with those fences being re-used for the next time around
vkQueueSubmit(graphicsQueue, 1, &sSubmitInfo, fence[nextIndex]);
```